



Version 0.8.0

- [1. Introduction](#)
- [2. Over-sampling](#)
- [3. Under-sampling](#)
- [4. Combination of over- and under-sampling](#)
- [5. Ensemble of samplers](#)**
- [6. Miscellaneous samplers](#)
- [7. Metrics](#)
- [8. Common pitfalls and recommended practices](#)
- [9. Dataset loading utilities](#)
- [10. Developer guideline](#)
- [11. References](#)

5. Ensemble of samplers

5.1. Classifier including inner balancing samplers

5.1.1. Bagging classifier

In ensemble classifiers, bagging methods build several estimators on different randomly selected subset of data. In scikit-learn, this classifier is named `BaggingClassifier`. However, this classifier does not allow to balance each subset of data. Therefore, when training on imbalanced data set, this classifier will favor the majority classes:

```
>>> from sklearn.datasets import make_classification
>>> X, y = make_classification(n_samples=10000,
...                           n_features=2, n_informative=2,
...                           n_redundant=0, n_repeated=0,
...                           n_classes=3,
...                           n_clusters_per_class=1,
...                           weights=[0.01, 0.05, 0.94],
...                           class_sep=0.8,
...                           random_state=0)
>>> from sklearn.model_selection import train_test_split
>>> from sklearn.metrics import balanced_accuracy_score
>>> from sklearn.ensemble import BaggingClassifier
>>> from sklearn.tree import DecisionTreeClassifier
>>> X_train, X_test, y_train, y_test = train_test_split(X,
...                                                    y,
...                                                    random_state=0)
>>> bc =
BaggingClassifier(base_estimator=DecisionTreeClassifier(),
...               random_state=0)
>>> bc.fit(X_train, y_train)
BaggingClassifier(...)
>>> y_pred = bc.predict(X_test)
>>> balanced_accuracy_score(y_test, y_pred)
0.77...
```

In `BalancedBaggingClassifier`, each bootstrap sample will be further resampled to achieve the `sampling_strategy` desired. Therefore, `BalancedBaggingClassifier` takes the same parameters than the scikit-learn `BaggingClassifier`. In addition, the sampling is controlled by the parameter `sampler` or the two parameters `sampling_strategy` and `replacement`, if one wants to use the `RandomUnderSampler`:

On this page

[5.1. Classifier including inner balancing samplers](#)

[Edit this page](#)

```
>>> from imblearn.ensemble import BalancedBaggingClassifier >>>
>>> bbc =
BalancedBaggingClassifier(base_estimator=DecisionTreeClassifier(
...
sampling_strategy='auto',
... replacement=False,
... random_state=0)
>>> bbc.fit(X_train, y_train)
BalancedBaggingClassifier(...)
>>> y_pred = bbc.predict(X_test)
>>> balanced_accuracy_score(y_test, y_pred)
0.8...
```

Changing the **sampler** will give rise to different known implementation [MO97], [HKT09], [WY09]. You can refer to the following example shows in practice these different methods: [Bagging classifiers using sampler](#)

5.1.2. Forest of randomized trees

[BalancedRandomForestClassifier](#) is another ensemble method in which each tree of the forest will be provided a balanced bootstrap sample [CLB+04]. This class provides all functionality of the [RandomForestClassifier](#):

```
>>> from imblearn.ensemble import
BalancedRandomForestClassifier >>>
>>> brf = BalancedRandomForestClassifier(n_estimators=100,
random_state=0)
>>> brf.fit(X_train, y_train)
BalancedRandomForestClassifier(...)
>>> y_pred = brf.predict(X_test)
>>> balanced_accuracy_score(y_test, y_pred)
0.8...
```

5.1.3. Boosting

Several methods taking advantage of boosting have been designed.

[RUSBoostClassifier](#) randomly under-sample the dataset before to perform a boosting iteration [SKVHN09]:

```
>>> from imblearn.ensemble import RUSBoostClassifier >>>
>>> rusboost = RUSBoostClassifier(n_estimators=200,
algorithm='SAMME.R',
... random_state=0)
>>> rusboost.fit(X_train, y_train)
RUSBoostClassifier(...)
>>> y_pred = rusboost.predict(X_test)
>>> balanced_accuracy_score(y_test, y_pred)
0...
```

A specific method which uses [AdaBoostClassifier](#) as learners in the bagging classifier is called “EasyEnsemble”. The [EasyEnsembleClassifier](#) allows to bag AdaBoost learners which are trained on balanced bootstrap

samples [LWZ08]. Similarly to the [BalancedBaggingClassifier](#) API, one can construct the ensemble as:

```
>>> from imblearn.ensemble import EasyEnsembleClassifier >>>
>>> eec = EasyEnsembleClassifier(random_state=0)
>>> eec.fit(X_train, y_train)
EasyEnsembleClassifier(...)
>>> y_pred = eec.predict(X_test)
>>> balanced_accuracy_score(y_test, y_pred)
0.6...
```

Examples

- [Compare ensemble classifiers using resampling](#)

[<< 4. Combination of over-
and under-sampling](#)

[6. Miscellaneous samplers >>](#)

© Copyright 2014-2021, The imbalanced-learn developers.

Created using [Sphinx](#) 3.4.3.